

ELK-day01

ELK-day01

- 1.ElasticSearch基本使用
- 2.ES单机安装:
- 3.ES索引基本操作
- 4.ES集群环境搭建
- 5.cerebro状态检查
- 6.集群节点
- 7.ES集群状态检查
- 8.ES集群扩展

ELK-day02

- 收集系统日志
- 收集Nginx
- 收集nginx访问日志和错误日志
- 收集nginx多个虚拟主机的日志
- Tomcat日志

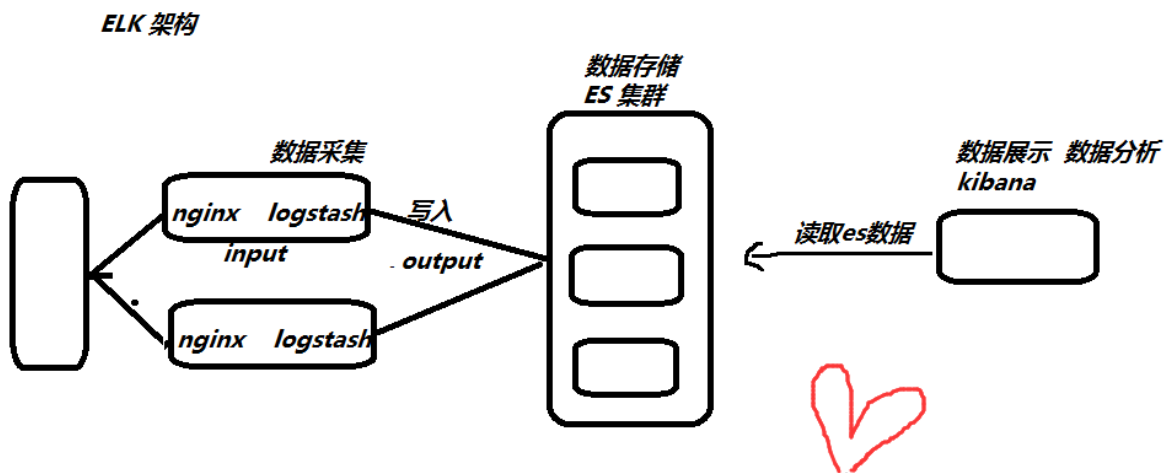
ELK-day03

- Logstash Input
- logstash 分析Nginx
- Logstash 分析 MySQL
- Logstash 分析 APP

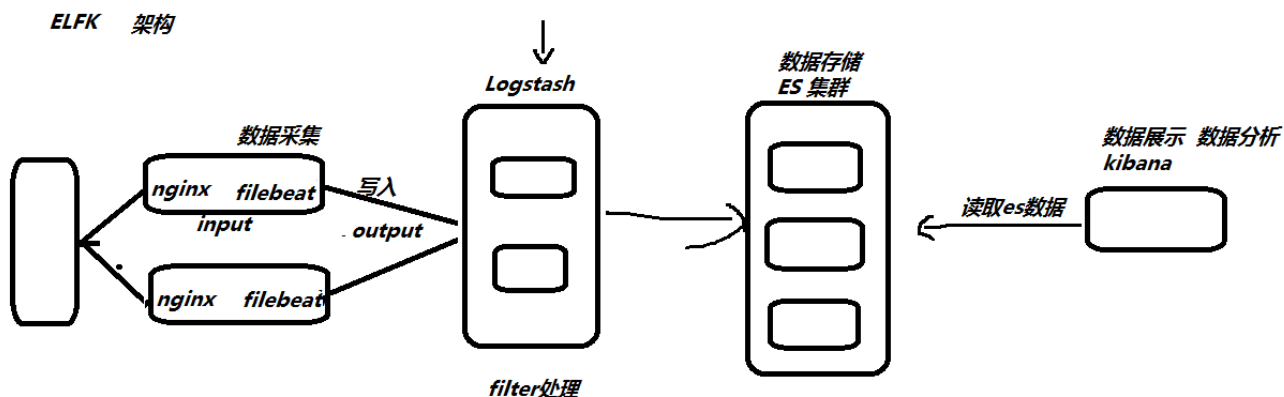
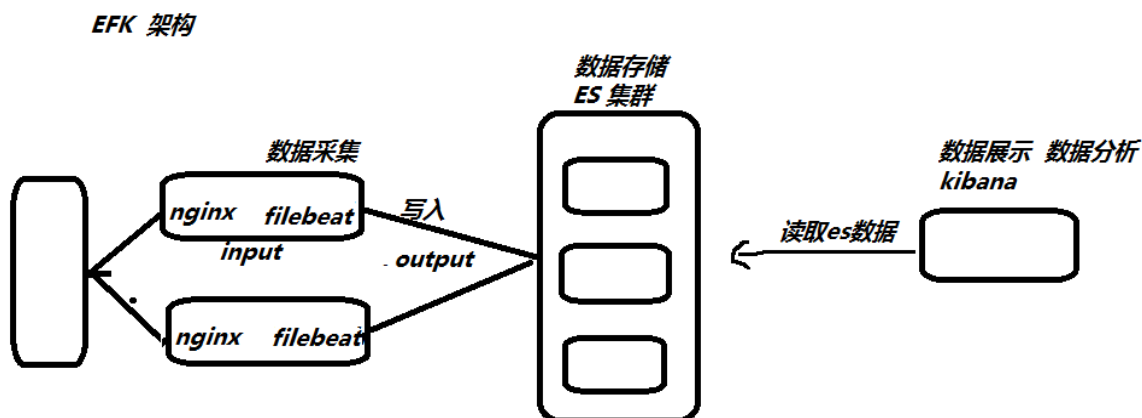
ELK-day04

- kibana
- Redis

ELK基础架构



E: elasticsearch 数据搜索 数据存储 java L: Logstash 数据收集 (数据解析 数据转换) 数据输出 java F: Filebeat 数据采集 (简单的数据处理) <--go K: Kibana 数据分析 数据展示



5.使用EFK收集哪些日志?

容器 : docker

代理 : Haproxy、Nginx

web : Nginx、Tomcat、Httpd、PHP

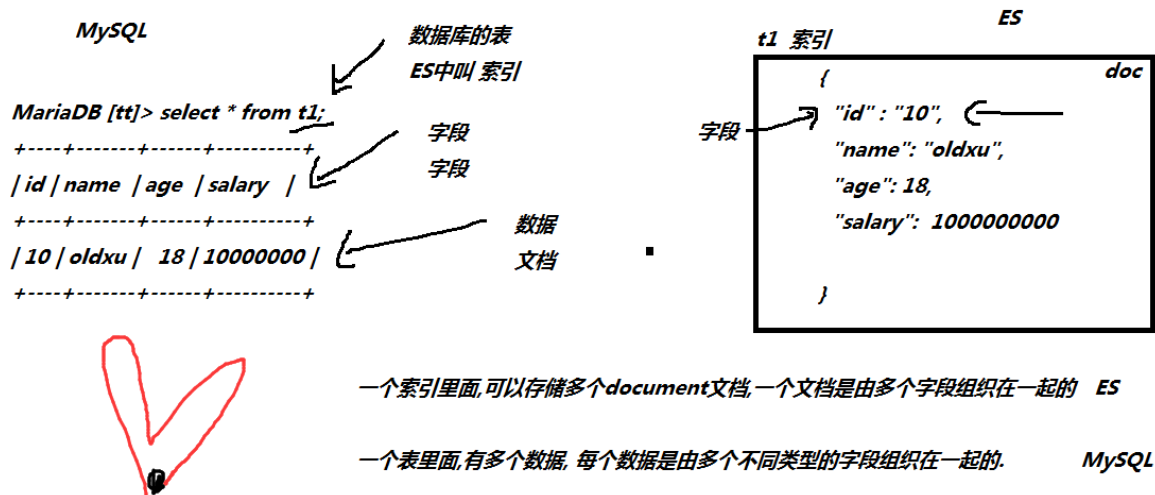
db : mysql、redis、mongo、elasticsearch

存储 : nfs、glusterfs、fastdfs

系统 : message、security

业务 : app

1.ElasticSearch基本使用



ES7

10.0.0.7 172.16.1.7 nginx+filebeat

10.0.0.8 172.16.1.8 nginx+filebeat

10.0.0.141 172.16.1.141 kafka

10.0.0.141 172.16.1.141 kafka

10.0.0.141 172.16.1.141 kafka

10.0.0.151 172.16.1.151 Logstash

10.0.0.152 172.16.1.152 Logstash

10.0.0.161 172.16.1.161 es-node1 ES Kibana

10.0.0.162 172.16.1.162 es-node2

10.0.0.163 172.16.1.163 es-node3

2.ES单机安装:

```
1 [root@es-node1 ~]# yum install java -y
2 [root@es-node1 ~]# rpm -ivh elasticsearch-7.4.0-
  x86_64.rpm kibana-7.4.0-x86_64.rpm
3
4 [root@es-node1 ~]# vim /etc/elasticsearch/jvm.options
5 -Xms512m      #实验环境 生产环境最少内存一半以上 官方建议 最高
  32Gb
6 -Xmx512m
7
8 [root@es-node1 ~]# systemctl enable
  elasticsearch.service
9 [root@es-node1 ~]# systemctl start
  elasticsearch.service
10
11 #测试es是否启动
12 [root@es-node1 ~]# curl 127.0.0.1:9200
```

```
13 {
14   "name" : "es-node1",
15   "cluster_name" : "elasticsearch",
16   "cluster_uuid" : "l8qVNdKlSvSqjUJSvUh4dw",
17   "version" : {
18     "number" : "7.4.0",
19     "build_flavor" : "default",
20     "build_type" : "rpm",
21     "build_hash" :
22       "22e1767283e61a198cb4db791ea66e3f11ab9910",
23     "build_date" : "2019-09-27T08:36:48.569419Z",
24     "build_snapshot" : false,
25     "lucene_version" : "8.2.0",
26     "minimum_wire_compatibility_version" : "6.8.0",
27     "minimum_index_compatibility_version" : "6.0.0-
beta1"
28   },
29   "tagline" : "You Know, for search"
30 }
```

```
31 #修改kibana的配置
```

```
32 [root@es-node1 ~]# vim /etc/kibana/kibana.yml
33 server.host: "0.0.0.0"
34 i18n.locale: "zh-CN"
```

```
35
```

```
36 #启动kibana
```

```
37 [root@es-node1 ~]# systemctl enable kibana
38 [root@es-node1 ~]# systemctl start kibana
```

```
39
```

```
40
```

```
41 #访问kibana
```

```
42
```

3.ES索引基本操作

创建一个索引

PUT /oldxu_es



查看所有的索引

GET _cat/indices

删除索引

DELETE /oldxu_es

1 # 创建一个索引

2 PUT /oldxu_es

3

4 # 查看所有的索引

5 GET _cat/indices

6

7 # 删除索引

8 DELETE /oldxu_es

9

10

11 #给oldxu_es索引录入一个文档

12 POST /tt/_doc/1

13 {

14 "name": "oldxu",

15 "age": 18,

16 "salary": 1000000000

17 }

18

19 POST /oldxu_es/_doc/2

```
20 {
21     "name": "oldguo",
22     "age": 35,
23     "salary": 100
24 }
25
26 #获取指定的id数据
27 GET /oldxu_es/_doc/1
28
29 #获取所有的文档 默认前10个
30 GET /oldxu_es/_search
31
32 #模糊查询
33 GET /oldxu_es/_search
34 {
35     "query": {
36         "term": {
37             "name": "oldxu"
38         }
39     }
40 }
41
42 #删除指定id的文档
43 DELETE /oldxu_es/_doc/1
44
45
46
47 #
48 POST _bulk
49 {"index":{"_index":"tt","_id":"1"}}
50 {"name":"oldxu","age":"18"}
51 {"create":{"_index":"tt","_id":"2"}}
52 {"name":"oldqiang","age":"30"}
53 {"delete":{"_index":"tt","_id":"2"}}
```

```
54 {"update":{"_id":"1","_index":"tt"}}
55 {"doc":{"age":"20"}}
56
57
58 #一次查询多个文档
59 GET _mget
60 {
61   "docs": [
62     {
63       "_index": "tt",
64       "_id": "1"
65     },
66     {
67       "_index": "tt",
68       "_id": "2"
69     }
70   ]
71 }
```

4.ES集群环境搭建

1 配置集群

```
1 #删除所有的es相关的数据（集群无法组件的情况）
2 [root@es-node1 ~]# rm -rf /var/lib/elasticsearch/*
3 [root@es-node1 ~]# systemctl stop elasticsearch.service
4 [root@es-node1 ~]# systemctl stop kibana
5
6 配置node1
7 [root@es-node1 ~]# grep '^[a-Z]'
   /etc/elasticsearch/elasticsearch.yml
8 cluster.name: my-olddxu
```



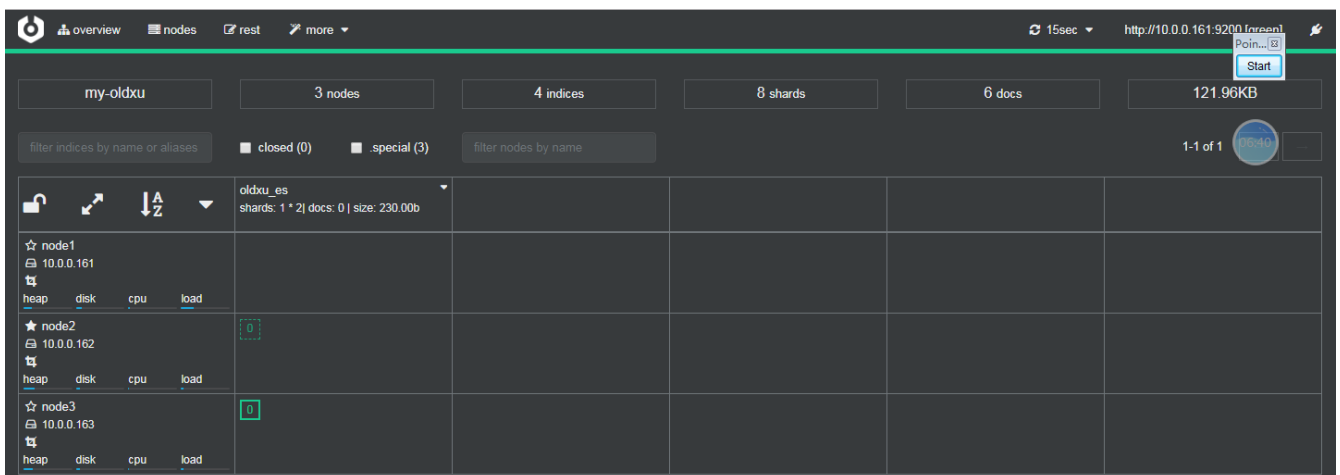
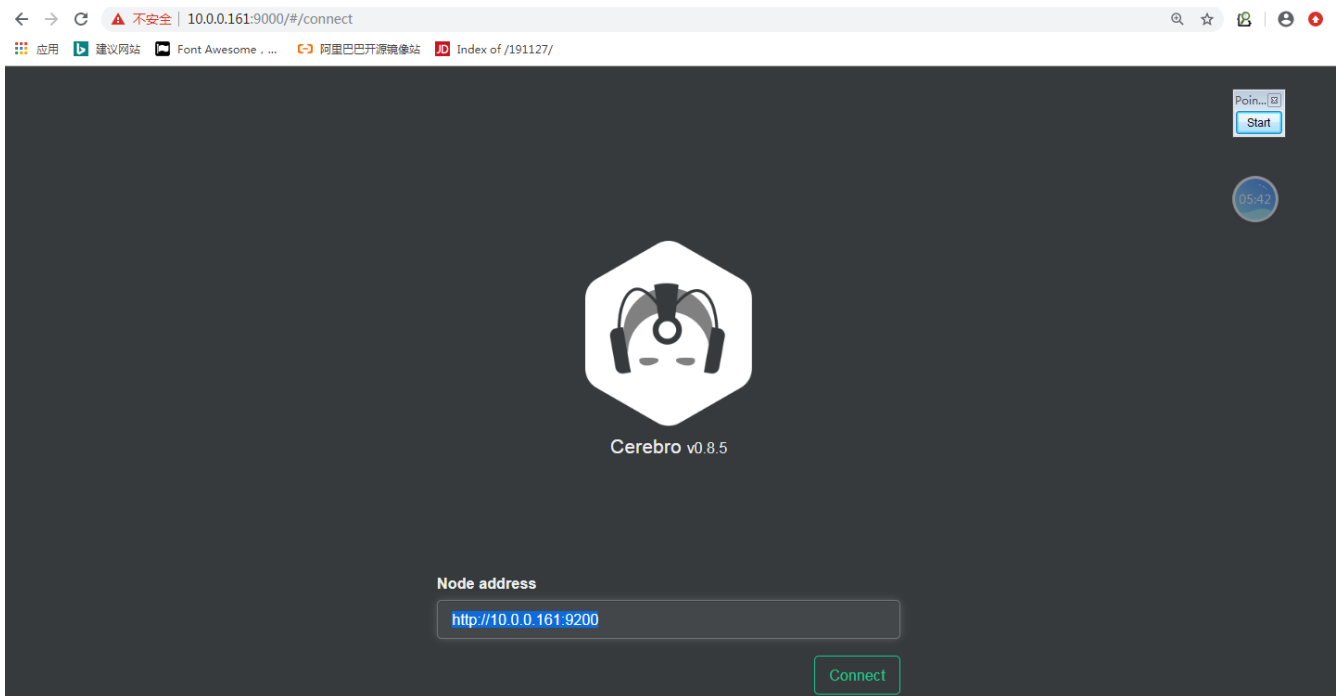
```
9 node.name: node1
10 path.data: /var/lib/elasticsearch
11 path.logs: /var/log/elasticsearch
12 network.host: 0.0.0.0
13 http.port: 9200
14 discovery.seed_hosts: ["10.0.0.161", "10.0.0.162",
    "10.0.0.163"]
15 cluster.initial_master_nodes: ["10.0.0.161",
    "10.0.0.162", "10.0.0.163"]
16
17 scp -rp /etc/elasticsearch/elasticsearch.yml
    root@172.16.1.162:/etc/elasticsearch/elasticsearch.yml
18 scp -rp /etc/elasticsearch/elasticsearch.yml
    root@172.16.1.163:/etc/elasticsearch/elasticsearch.yml
19
20 scp /etc/elasticsearch/jvm.options
    root@172.16.1.162:/etc/elasticsearch/jvm.options
21 scp /etc/elasticsearch/jvm.options
    root@172.16.1.163:/etc/elasticsearch/jvm.options
22
23
24 配置node2
25 [root@es-node2 ~]# grep "[a-z]"
    /etc/elasticsearch/elasticsearch.yml
26 cluster.name: my-olddxu
27 node.name: node2
28 path.data: /var/lib/elasticsearch
29 path.logs: /var/log/elasticsearch
30 network.host: 0.0.0.0
31 http.port: 9200
32 discovery.seed_hosts: ["10.0.0.161", "10.0.0.162",
    "10.0.0.163"]
33 cluster.initial_master_nodes: ["10.0.0.161",
    "10.0.0.162", "10.0.0.163"]
```

```
34
35
36 配置node3
37 [root@es-node3 ~]# grep "^[a-z]"
   /etc/elasticsearch/elasticsearch.yml
38 cluster.name: my-oldxu
39 node.name: node3
40 path.data: /var/lib/elasticsearch
41 path.logs: /var/log/elasticsearch
42 network.host: 0.0.0.0
43 http.port: 9200
44 discovery.seed_hosts: ["10.0.0.161", "10.0.0.162",
   "10.0.0.163"]
45 cluster.initial_master_nodes: ["10.0.0.161",
   "10.0.0.162", "10.0.0.163"]
46
47
48 启动所有节点
49 systemctl start elasticsearch
50
51 通过 curl 检查集群环境是否正常
52 curl http://10.0.0.163:9200/_cluster/health?pretty
53
54
55 kibana
56     GET /_cluster_health
57
58
```

5.cerebro状态检查

cerebro插件来检查整个集群的环境 默认监听9000端口

```
1 [root@es-node1 ~]# rpm -ivh cerebro-0.8.5-1.noarch.rpm
2 [root@es-node1 ~]# vim /etc/cerebro/application.conf
3 data.path = "/tmp/cerebro.db"
4 [root@es-node1 ~]# systemctl enable cerebro
```



6.集群节点

master角色: 负责控制整个集群的操作, 通过cluster_status状态维护集群.

选举: cluster.initial_master_nodes master-eligible

可以不参与选举: node.master: false cluster_state: 节点信息 索引信息

data角色: 存储数据 (默认都是data节点) 关闭data: node.data: false

coordinating角色: 负责路由 不能取消

7.ES集群状态检查

5.ES集群健康检查

Cluster Health 获取集群的健康状态, 整个集群状态包括以下三种:

- 1) green 健康状态, 指所有主副分片都正常分配
- 2) yellow 指所有主分片都正常分配, 但是有副本分片未正常分配
- 3) red 有主分片未分配, 表示索引不完备, 写也可能有问题。(但不代表不能存储数据和读取数据)
- 4) 可以通过 `GET _cluster/health?pretty=true` 方式获取集群状态

$10 \% 3 = 0$ 分片 $10 \% 3 = 1$ 分片 $10 \% 3 = 2$ 分片

```
1 shard = hash(routing) % number_of_primary_shards
2 # hash          算法保证将数据均匀分散在分片中
3 # routing       是一个关键参数, 默认是文档id, 也可以自定义。
4 # number_of_primary_shards 主分片数
5
6 # 注意: 该算法与主分片数相关, 一旦确定后便不能更改主分片。
7 # 因为一旦修改主分片修改后, shard的计算就完全不一样了。
```

8.ES集群扩展

ELK-day02



FILEBEAT

轻量型日志采集器

当您要面对成百上千、甚至成千上万的服务器、虚拟机和容器生成的日志时，请告别 SSH 吧。Filebeat 将为您提供一种轻量型方法，用于转发和汇总日志与文件，让简单的事情不再繁杂。

input 我们要采集的日志文件路径, 收割机 *harvester* 监听文件的变化 --> splooper程序 --> 转发 es | logstash | kafka | redis

filebeat配置-filebeat.yml

type类型有 log | stdin | redis | udp | docker

指定采集日志文件路径，可以使用*的方式

tags可以自定义标签

input可以采集多个日志文件

output输出至哪些存储
[elasticsearch| redis | logstash|kafka|.....

filebeat.inputs:

```
- type: log
  enabled: true
  paths:
    - /var/log/nginx/access.log
  tags: ["oldxu-access"]
```

```
- type: log
  enabled: true
  paths:
    - /var/log/nginx/error.log
  tags: ["oldxu-error"]
```

www.xuliangwei.com

```
output.elasticsearch:
  hosts: ["localhost:9092"]
```

```
1 filebeat.inputs:
2   - type: stdin          #标准输入
3     enabled: true        #启用
4
5 output.console:          #标准输出
6   pretty: true
7   enable: true
8
```

```
{
  "@timestamp": "2020-01-14T00:48:14.983Z",
  "@metadata": {
    "beat": "filebeat",
    "type": "_doc",
    "version": "7.4.0"
  },
  "log": {
    "offset": 0,
    "file": {
      "path": ""
    }
  },
}
```

```
1 #将文件最新发生变化的内容,存入ES
2
3 [root@web01 ~]# cat /etc/filebeat/file.yml
4 filebeat.inputs:
5   - type: log
6     paths: /var/log/nginx/access.log
7     enabled: true
8
9 output.elasticsearch:
10   hosts:
      ["10.0.0.161:9200","10.0.0.162:9200","10.0.0.163:9200"]
```

收集系统日志

特别分散--> syslog --> file.txt

1.减少无用的数据

2.调整索引名称

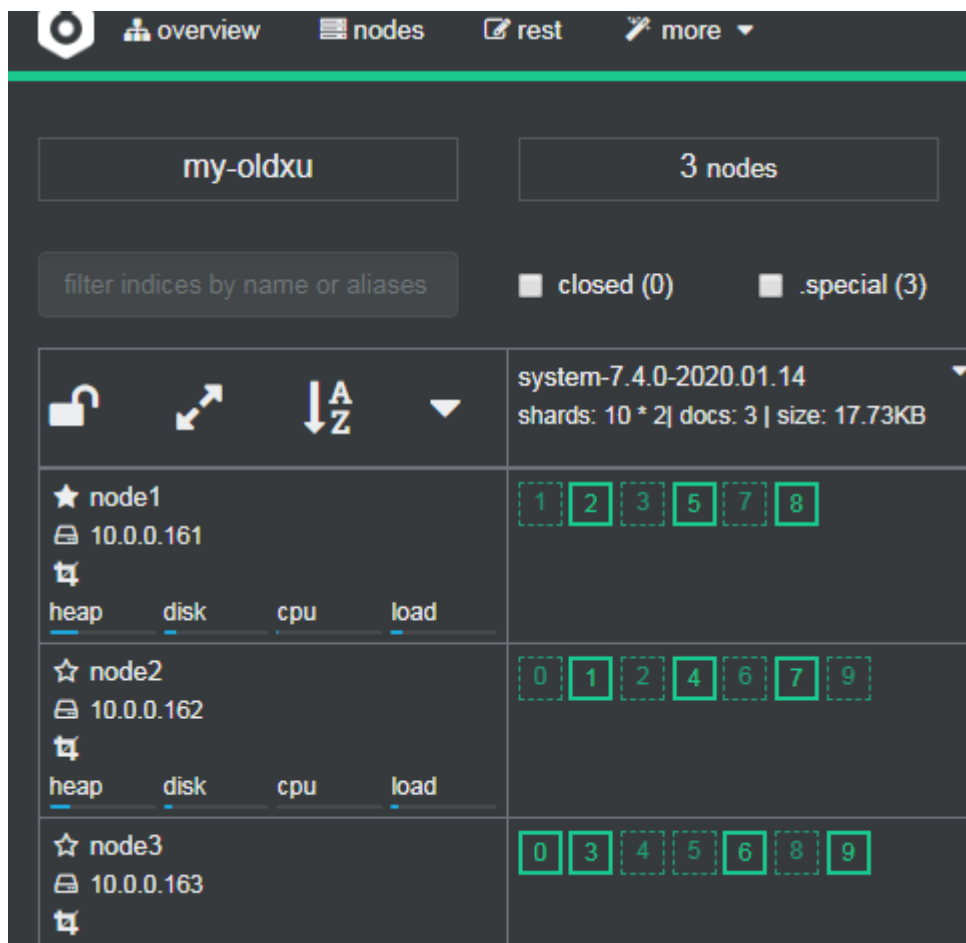
3.测试调整模板,设定分片

```
1 [root@web01 ~]# cat /etc/filebeat/filebeat.yml
2 filebeat.inputs:
3   - type: log
4     enabled: true
5     paths:
6       - /var/log/oldxu.log
7     include_lines: ['^ERR', '^WARN', 'sshd']    #只看指定的
      日志
8
9 output.elasticsearch:
10   hosts:
11     ["10.0.0.161:9200","10.0.0.162:9200","10.0.0.163:9200"]
12   index: "system-%{[agent.version]}-%{+yyyy.MM.dd}"
13
14 setup.ilm.enabled: false
15 setup.template.name: system    #索引关联的模板名称
16 setup.template.pattern: system-*
17
18
19 方式一:
```

```

20 ###设定system模板的分片数和副本数
21 #setup.template.settings:                #定义索引分片数和副本
22 #  index.number_of_shards: 3
23 #  index.number_of_replicas: 1
24
25 方式二:
26
27     "number_of_routing_shards": "30",
28     "number_of_shards": "10",
29     "number_of_replicas": "1",
30
31     1.修改system模板    --->   添加 shards 分片数数
    量,replicas的数量
32     2.删除模板关联的索引
33     3.删除filebeat自行指定的分片数和副本数
34     4.重启filebeat
35     5.产生新的日志

```



收集Nginx

```
1 log_format json '{ "time_local": "$time_local", '
2                   '"remote_addr":
3                   '"$remote_addr", '
4                   '"referer": "$http_referer",
5                   '"request": "$request", '
6                   '"status": $status, '
7                   '"bytes": $body_bytes_sent, '
8                   '"agent": "$http_user_agent",
9                   '"x_forwarded":
10                  '"$http_x_forwarded_for", '
11                  '"up_addr":
12                  '"$upstream_addr", '
13                  '"up_host":
14                  '"$upstream_http_host", '
15                  '"upstream_time":
16                  '"$upstream_response_time", '
17                  '"request_time":
18                  '"$request_time"'
19                '}' ;
20
21 access_log /var/log/nginx/access.log json;
```

配置filebeat

```
1 [root@web01 filebeat]# cat /etc/filebeat/filebeat.yml
```

```

2 filebeat.inputs:
3   - type: log
4     enabled: true
5     paths:
6       - /var/log/nginx/access.log
7     json.keys_under_root: true      #默认false, 还会将json解析
    的日志存储至messages字段
8     json.overwrite_keys: true      #覆盖默认的关键字, 使用自定义
    json格式的关键字
9
10  output.elasticsearch:
11    hosts:
12      ["10.0.0.161:9200", "10.0.0.162:9200", "10.0.0.163:9200"]
13    index: "nginx-%{[agent.version]}-%{+yyyy.MM.dd}"
14
15  setup.ilm.enabled: false
16  setup.template.name: nginx      #索引关联的模板名称
17  setup.template.pattern: nginx-*

```

收集nginx访问日志和错误日志

```

1 [root@web01 filebeat]# cat filebeat.yml
2 filebeat.inputs:
3   - type: log
4     enabled: true
5     paths:
6       - /var/log/nginx/access.log
7     json.keys_under_root: true      #默认false, 还会将json解析
    的日志存储至messages字段
8     json.overwrite_keys: true      #覆盖默认的关键字, 使用自定义
    json格式的关键字
9     tags: ["access"]
10
11

```

```
12 - type: log
13   enabled: true
14   paths:
15     - /var/log/nginx/error.log
16   tags: ["error"]
17
18
19
20 output.elasticsearch:
21   hosts:
22     ["10.0.0.161:9200", "10.0.0.162:9200", "10.0.0.163:9200"]
23   indices:
24     - index: "nginx-access-%{[agent.version]}-%
25       {+yyyy.MM.dd}"
26       when.contains:
27         tags: "access"
28
29     - index: "nginx-error-%{[agent.version]}-%
30       {+yyyy.MM.dd}"
31       when.contains:
32         tags: "error"
33
34 setup.ilm.enabled: false
35 setup.template.name: nginx    #索引关联的模板名称
36 setup.template.pattern: nginx-*
```


收集nginx多个虚拟主机的日志

elk.oldxu.com

bk.oldxu.com

bs.oldxu.com

error日志



1.虚拟主机

```
1 [root@web01 conf.d]# cat elk.oldxu.com.conf
2 server {
3     listen 80;
4     server_name elk.oldxu.com;
5     root /code/elk;
6     access_log /var/log/nginx/elk.oldxu.com.log json;
7
8     location / {
9         index index.html;
10    }
11 }
12
13
14 [root@web01 conf.d]# cat bs.oldxu.com.conf
15 server {
16     listen 80;
17     server_name bs.oldxu.com;
```

```
18     root /code/bs;
19     access_log /var/log/nginx/bs.olddxu.com.log json;
20
21     location / {
22         index index.html;
23     }
24 }
25
26 [root@web01 conf.d]# cat bk.olddxu.com.conf
27 server {
28     listen 80;
29     server_name bk.olddxu.com;
30     root /code/bk;
31     access_log /var/log/nginx/bk.olddxu.com.log json;
32
33     location / {
34         index index.html;
35     }
36 }
```

2.测试,模拟产生日志

```
1 [root@web01 conf.d]# curl -H Host:elk.olddxu.com
  http://10.0.0.7
2 elk.olddux.com
3 [root@web01 conf.d]# curl -H Host:bs.olddxu.com
  http://10.0.0.7
4 bs.olddux.com
5 [root@web01 conf.d]# curl -H Host:bk.olddxu.com
  http://10.0.0.7
6 bk.olddux.com
7
```

3.配置filebeat

```
1 [root@web01 filebeat]# cat /etc/filebeat/filebeat.yml
2 filebeat.inputs:
3 - type: log
4   enabled: true
5   paths:
6     - /var/log/nginx/elk.olddxu.com.log
7   json.keys_under_root: true
8   json.overwrite_keys: true
9   tags: ["nginx-elk-host"]
10
11 - type: log
12   enabled: true
13   paths:
14     - /var/log/nginx/bs.olddxu.com.log
15   json.keys_under_root: true
16   json.overwrite_keys: true
17   tags: ["nginx-bs-host"]
18
19 - type: log
20   enabled: true
21   paths:
22     - /var/log/nginx/bk.olddxu.com.log
23   json.keys_under_root: true
24   json.overwrite_keys: true
25   tags: ["nginx-bk-host"]
26
27
28 - type: log
29   enabled: true
30   paths:
31     - /var/log/nginx/error.log
32   tags: ["nginx-error"]
33
34
```

```
35 output.elasticsearch:
36   hosts:
37     ["10.0.0.161:9200","10.0.0.162:9200","10.0.0.163:9200"]
37   indices:
38     - index: "nginx-elk-access-%{[agent.version]}-%
39       {+yyyy.MM.dd}"
40       when.contains:
41         tags: "nginx-elk-host"
42     - index: "nginx-bs-access-%{[agent.version]}-%
43       {+yyyy.MM.dd}"
44       when.contains:
45         tags: "nginx-bs-host"
46     - index: "nginx-bk-access-%{[agent.version]}-%
47       {+yyyy.MM.dd}"
48       when.contains:
49         tags: "nginx-bk-host"
50     - index: "nginx-error-%{[agent.version]}-%
51       {+yyyy.MM.dd}"
52       when.contains:
53         tags: "nginx-error"
54 setup.ilm.enabled: false
55 setup.template.name: nginx    #索引关联的模板名称
56 setup.template.pattern: nginx-*
```

Tomcat日志

访问日志 ---> json格式

```

1 #1.修改tomcat日志格式
2
3 [root@web02 soft]# yum install java -y
4 [root@web02 soft]# vim tomcat/conf/server.xml
5     <Host name="tomcat.olddxu.com" appBase="webapps"
6         unpackWARs="true" autoDeploy="true">
7         <Valve
8             className="org.apache.catalina.valves.AccessLogValve"
9             directory="logs"
10                prefix="tomcat.olddxu.com.log"
11                suffix=".txt"
12                pattern="
13                    {&quot;clientip&quot;:&quot;%h&quot;,&quot;ClientUser&
14                    uot;:&quot;%l&quot;,&quot;authenticated&quot;:&quot;%u&
15                    quot;,&quot;AccessTime&quot;:&quot;%t&quot;,&quot;metho
16                    d&quot;:&quot;%r&quot;,&quot;status&quot;:&quot;%s&quot;
17                    ;,&quot;SendBytes&quot;:&quot;%b&quot;,&quot;Query?
18                    string&quot;:&quot;%q&quot;,&quot;partner&quot;:&quot;%
19                    {Referer}i&quot;,&quot;AgentVersion&quot;:&quot;%{User-
20                    Agent}i&quot;}" />
21    </Host>

```

配置filebeat

```

1 [root@web02 filebeat]# cat /etc/filebeat/filebeat.yml
2 filebeat.inputs:
3 - type: log
4   enabled: true
5   paths:
6     - /soft/tomcat/logs/tomcat.olddxu.com.log.*.txt
7   json.keys_under_root: true      #默认False, 还会将json解析
  的日志存储至messages字段
8   json.overwrite_keys: true      #覆盖默认的关键字, 使用自定义
  json格式的关键字

```



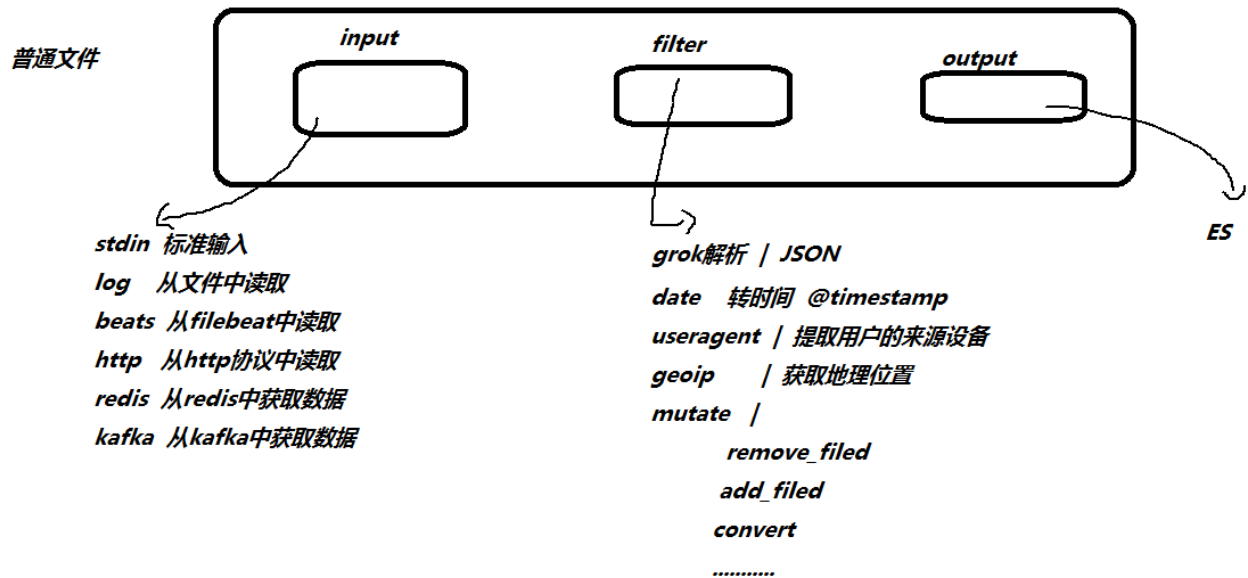
```
9
10
11 output.elasticsearch:
12   hosts: ["10.0.0.161:9200","10.0.0.162:9200"]
13   index: "tomcat-access-%{[agent.version]}-%
14     {+yyyy.MM.dd}"
15 setup.ilm.enabled: false
16 setup.template.name: tomcat    #索引关联的模板名称
17 setup.template.pattern: tomcat-*
18
```

错误日志 <--java

```
1 [root@web02 filebeat]# cat filebeat.yml
2 filebeat.inputs:
3 - type: log
4   enabled: true
5   paths:
6     - /soft/tomcat/logs/tomcat.olddxu.com.log.*.txt
7   json.keys_under_root: true    #默认false, 还会将json解析
  的日志存储至messages字段
8   json.overwrite_keys: true    #覆盖默认的key, 使用自定义
  json格式的key
9   tags: ["tomcat-access"]
10
11 - type: log
12   enabled: true
13   paths:
14     - /soft/tomcat/logs/catalina.out
15   multiline.pattern: '^\\d{2}'  #匹配以2个数字开头的
16   multiline.negate: true
17   multiline.match: after
```

```
18 multiline.max_lines: 10000      #默认最大合并行为500，可根据
   实际情况调整。
19 tags: ["tomcat-error"]
20
21
22 output.elasticsearch:
23   hosts: ["10.0.0.161:9200","10.0.0.162:9200"]
24   indices:
25     - index: "tomcat-access-%{[agent.version]}-%
      {+yyyy.MM.dd}"
26       when.contains:
27         tags: "tomcat-access"
28
29     - index: "tomcat-error-%{[agent.version]}-%
      {+yyyy.MM.dd}"
30       when.contains:
31         tags: "tomcat-error"
32
33
34 setup.ilm.enabled: false
35 setup.template.name: tomcat      #索引关联的模板名称
36 setup.template.pattern: tomcat-*
```

ELK-day03



Logstash Input

```

1
2 [root@logstash-node1 conf.d]# cat
  input_file_output_console.conf
3 input {
4     file {
5         path => "/var/log/oldxu.log"
6         type => syslog
7         exclude => "*.gz"           #不想监听的文件规则，基于
glob匹配语法
8         start_position => "beginning"   #第一次从头开始读
取文件 beginning or end
9         stat_interval => "3"           #定时检查文件是否更新，默认
1s
10     }
11 }
12
13 output {
14     stdout {
15         codec => rubydebug

```

```

16     }
17 }
18
19
20
21 [root@logstash-node1 conf.d]# cat
   input_stdin_output_console.conf
22 input {
23     stdin {
24         type => stdin
25         tags => "tags_stdin"
26     }
27
28 }
29 output {
30     stdout {
31         codec => "rubydebug"
32     }
33 }

```

logstash 分析Nginx

```

1 66.249.73.135 - - [20/May/2015:21:05:11 +0000] "GET
   /blog/tags/xsendevent HTTP/1.1" 200 10049 "-"
   "Mozilla/5.0 (iPhone; CPU iPhone OS 6_0 like Mac OS x)
   AppleWebKit/536.26 (KHTML, like Gecko) version/6.0
   Mobile/10A5376e Safari/8536.25 (compatible;
   Googlebot/2.1; +http://www.google.com/bot.html)"

```

```

1 [root@logstash-node1 conf.d]# cat
   input_filebeat_output_es.conf
2 input {

```

```
3     beats {
4         port => 5044
5     }
6 }
7
8 filter {
9
10 if "nginx-access" in [tags][0] {
11     grok {
12         match => { "message" => "%{IPORHOST:clientip} %{
13             {USER:ident} %{USER:auth} \[%{HTTPDATE:timestamp}\] \"
14             (?:%{WORD:verb} %{NOTSPACE:request}(?: HTTP/%
15             {NUMBER:httpversion})?|%{DATA:rawrequest})\" %
16             {NUMBER:response} (?:%{NUMBER:bytes}|-) %{QS:referrer}
17             %{QS:useragent}\" }
18         }
19
20     date {
21         match => ["timestamp", "dd/MMM/yyyy:HH:mm:ss
22             z"]
23
24         target => "@timestamp"
25         timezone => "Asia/Shanghai"
26     }
27
28     geoip {
29         source => "clientip"
30     }
31
32     useragent {
33         source => "useragent"
34         target => "useragent"
35     }
36
37     mutate {
```

```

31         rename => ["%{[host][name]}" , "hostname" ]
32         convert => [ "bytes", "integer" ]
33         remove_field => [ "message", "agent" ,
"input","ecs" ]
34         add_field => { "target_index" => "logstash-
nginx-access-%{+YYYY.MM.dd}" }
35     }
36 } else if "nginx-error" in [tags][0] {
37     mutate {
38         add_field => { "target_index" => "logstash-
nginx-error-%{+YYYY.MM.dd}" }
39     }
40 }
41
42 }
43
44 output {
45     elasticsearch {
46         hosts =>
["10.0.0.161:9200","10.0.0.162:9200","10.0.0.163:9200"]
47         index => "%{[target_index]}"
48     }
49 }
50

```

Logstash 分析 MySQL

filebeat

```

1 [root@web01 filebeat]# cat filebeat.yml
2 filebeat.inputs:

```

```
3 - type: log
4   enabled: true
5   paths:
6     - /var/log/mariadb/slow.log
7   exclude_lines: ['^\# Time']
8   multiline.pattern: '^\# User'
9   multiline.negate: true
10  multiline.match: after
11  multiline.max_lines: 10000
12  tags: ["mysql-slow"]
13
14 output.logstash:
15   hosts: ["10.0.0.151:5044"]
```

logstash

```
1 [root@logstash-node1 conf.d]# cat
  input_filebeat_mysql_output_es.conf
2 input {
3     beats {
4         port => 5044
5     }
6 }
7
8
9 filter {
10
11     mutate {
12         gsub => ["message","\n"," "]
13     }
14     grok {
15
16         match => {
```

```

17     "message" => "(?m)^# User@Host: %{USER:User}\[%
{USER-2:User}\] @ (?: (?<Clienthost>\S*) )? \[ (?: %
{IP:Client_IP})? \] # Thread_id: %
{NUMBER:Thread_id:integer} \s+ Schema: (?: (?<DBname>\S*)
)\s+ QC_hit: (?: (?<QC_hit>\S*) ) # Query_time: %
{NUMBER:Query_Time} \s+ Lock_time: %
{NUMBER:Lock_Time} \s+ Rows_sent: %
{NUMBER:Rows_Sent:integer} \s+ Rows_examined: %
{NUMBER:Rows_Examined:integer} SET timestamp=%
{NUMBER:timestamp}; \s*(?<Query>( ?<Action>\w+)\s+.*)"
18     }
19 }
20
21 date {
22     match => ["timestamp", "UNIX", "YYYY-MM-dd
HH:mm:ss"]
23     target => "@timestamp"
24     timezone => "Asia/Shanghai"
25 }
26 mutate {
27     remove_field =>
["message", "input", "timestamp", "agent", "ecs", "log"]
28     convert => ["Lock_Time", "float"]
29     convert => ["Query_Time", "float"]
30     add_field => { "target_index" => "logstash-
mysql-slow-%{+YYYY.MM.dd}" }
31 }
32 }
33
34 output {
35     elasticsearch {
36         hosts => ["10.0.0.161:9200"]
37         index => "%{[target_index]}"
38     }

```



```
39     stdout {
40         codec => "rubydebug"
41     }
42 }
```

Logstash 分析 APP

模拟产生日志

```
1 | java -jar app-dashboard-1.0-SNAPSHOT.jar
   &>/var/log/app.log
```

filebeat

```
1 filebeat.inputs:
2 - type: log
3   enabled: true
4   paths:
5     - /var/log/app.log
6
7 output.logstash:
8   hosts: ["10.0.0.151:5044"]
```

logstash

```
1 [root@logstash-node1 conf.d]# cat
  input_filebeat_app_output_es.conf
2 input {
3     beats {
4         port => 5044
5     }
6 }
7
8 filter {
```

```
9      mutate {
10          split => {"message" => "|"}
11          add_field => {
12              "UserID" => "%{[message][1]}"
13              "Action" => "%{[message][2]}"
14              "Date" => "%{[message][3]}"
15          }
16
17          convert => {
18              "UserID" => "integer"
19              "Action" => "string"
20              "Date"    => "string"
21          }
22      }
23
24      #2020-01-15 17:04:15
25      date {
26          match => ["Date", "yyyy-MM-dd HH:mm:ss"]
27          target => "@timestamp"
28          timezone => "Asia/Chongqing"
29      }
30
31      mutate {
32          #remove_field => ["message", "Date"]
33          add_field => { "target_index" => "logstash-app-
34      %{+YYYY.MM.dd}" }
35      }
36
37      output {
38          elasticsearch {
39              hosts => ["10.0.0.161:9200"]
40              index => "%{[target_index]}"
41              template_overwrite => true
```

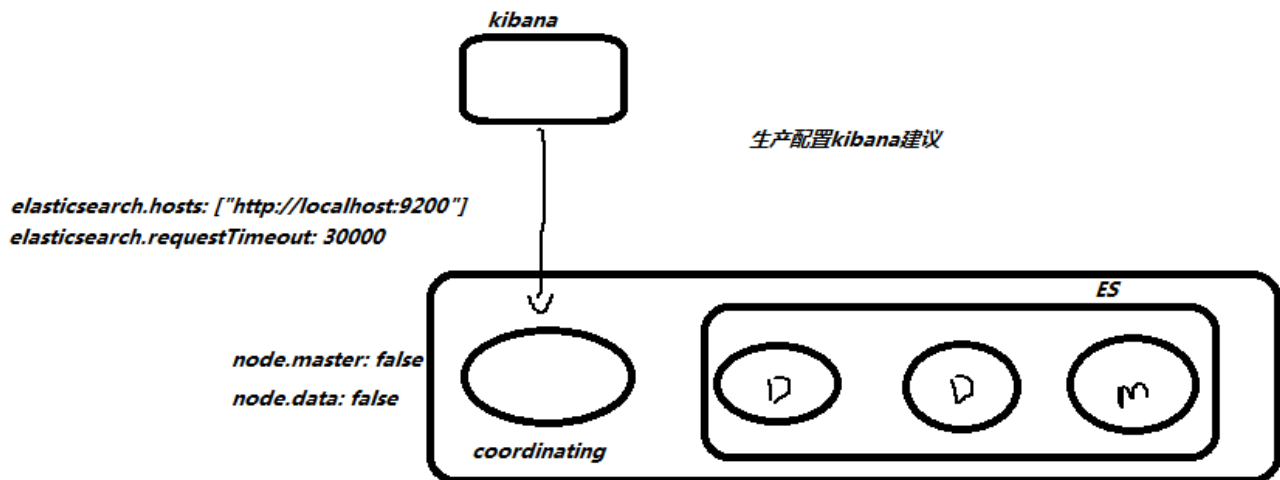
```
42     }
43 }
```

ELK-day04

```
1  -----if "access" in [tags][0]                                else
   if
2  input
3      beats --> redis|kafka(解耦|缓存) --> logstash (数据
   解析|数据转换) --> ES
4      http
5      file
6      redis
7      kafka
8  filter
9      grok    <--一推正则表达式的别名    | 将非结构化数据转为结构化
   数据
10     useragent    mac | linux | windows    小米    iphone
11     geoip        城市分布 | 经纬度
12     date         时间转换    时区
13     mutate
14         add_field
15         remove_field
16         rename_field
17         split
18         gsub
19         convert
20 output
21     ES
```

kibana

kibana用来做数据的展示和数据的分析,读取ES的索引|获取对应的数据



- 统计网站总PV、IP
- 统计网站每天PV、IP
- 统计访问IP Top10
- 统计来源的referrer
- 统计前10的资源 标签云
- 统计访问状态码
- 统计客户端设备
- 统计大于1s请求 待会出
- 统计网站总流量
- 统计网站访问流量趋势图

```
1 [root@web01 bk]# cat /etc/filebeat/filebeat.yml
2 filebeat.inputs:
3 - type: log
4   enabled: true
5   paths:
6     - /var/log/nginx/access.log
7   tags: ["access"]
```

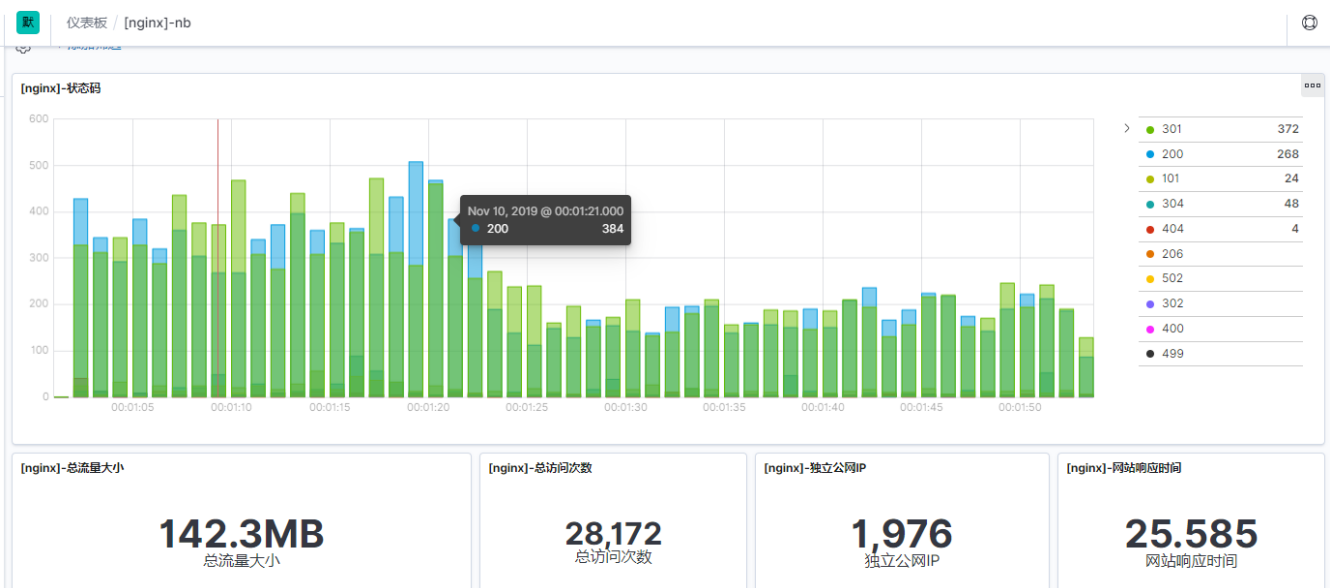
```
8
9
10 #- type: log
11 #   enabled: true
12 #   paths:
13 #     - /var/log/nginx/error.log
14 #   tags: ["error"]
15
16
17 output.logstash:
18   hosts: ["10.0.0.151:5044"]
19
20
21
22
23 tash-node1 conf.d]# cat
   input_filebeat_nginx_output_es.conf
24
25 input {
26     beats {
27         port => 5044
28     }
29 }
30
31 filter {
32     grok {
33         match => { "message" => "%
{COMBINEDAPACHELOG}" }
34     }
35
36     geoip {
37         source => "clientip"
38     }
39
```






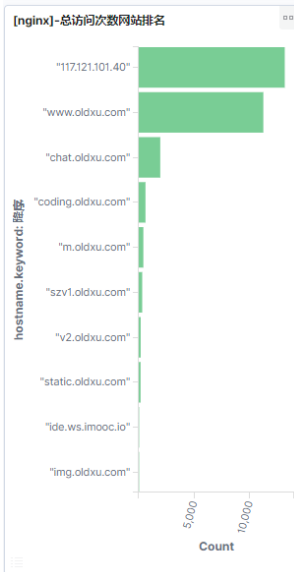
```

40     #30/Dec/2019:11:59:18 +0800
41     date {
42         match => ["timestamp", "dd/MMM/yyyy:HH:mm:ss
z"]
43         target => "@timestamp"
44         timezone => "Asia/Shanghai"
45     }
46
47     useragent {
48         source => "agent"
49         target => "agent"
50     }
51
52     mutate {
53         remove_field => [ "message","timestamp" ]
54         convert => [ "bytes" , "integer" ]
55     }
56
57 }
58
59 output {
60     stdout {
61         codec => rubydebug
62     }
63
64
65     elasticsearch {
66         hosts =>
["10.0.0.161:9200","10.0.0.162:9200","10.0.0.163:9200
"]
67         index => "kibana-nginx-%{+YYYY.MM.dd}"      #
索引名称
68         template_overwrite => true
69     }

```

```
1 #www.oidxu.com 为请求的域名 api3为请求的资源 uid=为用户id
2 124.161.176.119 - - [10/Nov/2019:00:01:52 +0800] "POST
  /api3/appover HTTP/1.1" 200 103 "www.oidxu.com" "-"
  code=B157963E-5BDA-4090-A021-
  A3D46D2E6BA2&secret=f0fbb455c7aebc69c5cc39d68c7859fe&
  time=9441&timestamp=1478707307012&token=12d0f0cfa1efb8
  1e42c321f027bbe752&uid=4384521 "oidxu/5.0.2 (iPhone;
  iOS 10.1.1; scale/2.00)" "-" 10.100.136.65:80 200
  0.011 0.011
```






[nginx]-访问来源

www.baidu.com
www.oidxu.com
coding.oidxu.com
m.oidxu.com
blog.jobbole.com


[nginx]-地域分布




[nginx]-浏览器设备



[nginx]-终端设备



[nginx]-操作系统分布



nginx-discover

1-50 of 28172

Time	_source
Nov 10, 2019 @ 00:01:53.000	<pre>ident: - @version: 1 referer_host: www.oidxu.com bytes: 538 target_index: logstash-kibana-nginx-access-2019.11 request: /course/ajaxcourselearn?cid=422 useragent_name: Chrome useragent.minor: 0 useragent.major: 45 useragent.device: Other useragent.os: Windows useragent.patch: 2454 useragent.os_name: Windows useragent.build: oidxu_type: learn x_forward_for: "-" method: GET response_time: 0.326 post_args: - upstream_host: 10.100.136.64:80 port: 80 @timestamp: Nov 10, 2019 @ 00:01:53.000 timestamp: 10/Nov/2019:00:01:53 +0800 response: 200 geolp.country_name: China geolp.location: ("lon": 118.7778, "lat": 32.0617) geolp.timezone: Asia/Shanghai geolp.country_code3: CN geolp.city_name: Nanjing</pre>

Redis

收集容器日志

<https://www.cnblogs.com/xuliangwei/p/12172960.html>